

Breast Cancer Prediction Project

HarvardX: PH125.9x Data Science - Choose your own project

Gabriele Mineo

17 April, 2019

Contents

1	Overview	3
1.1	Introduction	3
1.2	Aim of the project	4
1.3	Dataset	4
2	Methods and Analysis	6
2.1	Data Analysis	6
2.2	Modelling Approach	14
2.2.1	Modelling	14
2.2.2	Model creation	19
2.2.3	Naive Bayes Model	19
2.2.4	Logistic Regression Model	21
2.2.5	Random Forest Model	23
2.2.6	K Nearest Neighbor (KNN) Model	24
2.2.7	Neural Network with PCA Model	26
2.2.8	Neural Network with LDA Model	27
3	Results	29
4	Discussion	32
5	Conclusion	34
6	Appendix - Environment	35

Chapter 1

Overview

This project is related to the Choose-your-own project of the HarvardX: PH125.9x Data Science: Capstone course. The present report starts with a general idea of the project and by representing its objectifs.

Then the given dataset will be prepared and setup. An exploratory data analysis is carried out in order to develop a machine learning algorithm that could predict whether a breast cancer cell is benign or malignant until a final model. Results will be explained. Finally, the report will end with some concluding remarks.

1.1 Introduction

A neoplasm is an abnormal mass of tissue, the growth of which exceeds and is uncoordinated with that of the normal tissues, and persists in the same excessive manner after cessation of the stimulus which evoked the change. Cancer can start almost anywhere in the human body, which is made up of 37.200 billion cells. As these tumors grow, some cancer cells can break off and travel to distant places in the body through the blood or the lymph system and form new tumors far from the original one. Unlike malignant tumors, benign tumors do not spread into, or invade, nearby tissues. Breast cancer refers to a pathology in which a tumor develops in the breast tissue. Breast cancer is amongst the most common type of cancer in both sexes since 1975 and causes around 411,000 annual deaths worldwide. It is predicted that the incidence for worldwide cancer will continue to increase, with 23,6 million new cancer cases each year by 2030, corresponding to 68% more cases in comparison to 2012.

Mammography is the most common mass screening tool for an early detection of breast cancers because of its sensitivity in recognising breast masses. After detection of suspicious breast masses, a biopsy test procedure would be carried out, such as Fine Needle Aspirates (FNA), that is the method this report focus on. This method has been showed to be safe, cost-effective, accurate and fast. A small drop of viscous fluid is aspired from the breast masses to be analysed under the microscope. Then, a small region of the breast mass cells is photographed in a grey scale image and further analysed using an image analysis program 'Xcyt'. This program uses a curve-fitting to determine the edges of the nuclei from initial dots manually placed near these edges by a mouse.

The edges of the visible cell nuclei were manually placed with a mouse (red dots), 'Xcyt' program will after outline the nuclei (red circle). The interactive diagnosis process takes about 5 minutes per sample.

This project will make a performance comparison between different machine learning algorithms in order to assess the correctness in classifying data with respect to efficiency and effectiveness of each algorithm in terms of accuracy, precision, sensitivity and specificity, in order to find the best diagnosis.

Diagnosis in an early stage is essential to the facilitate the subsequent clinical management of patients and increase the survival rate of breast cancer patients.

The major models used and tested will be supervised learning models (algorithms that learn from labelled data), which are most used in these kinds of data analysis.

The utilization of data science and machine learning approaches in medical fields proves to be prolific as such approaches may be considered of great assistance in the decision making process of medical practitioners. With an unfortunate increasing trend of breast cancer cases, comes also a big deal of data which is of significant use in furthering clinical and medical research, and much more to the application of data science and machine learning in the aforementioned domain.

1.2 Aim of the project

The objective of this report is to train machine learning models to predict whether a breast cancer cell is Benign or Malignant. Data will be transformed and its dimension reduced to reveal patterns in the dataset and create a more robust analysis. As previously said, the optimal model will be selected following the resulting accuracy, sensitivity, and f1 score, amongst other factors. We will later define these metrics. We can use machine learning method to extract the features of cancer cell nuclei image and classify them. It would be helpful to determine whether a given sample appears to be Benign (“B”) or Malignant (“M”).

The machine learning models that we will applicate in this report try to create a classifier that provides a high accuracy level combined with a low rate of false-negatives (high sensitivity).

1.3 Dataset

The present report covers the Breast Cancer Wisconsin (Diagnostic) DataSet (<https://www.kaggle.com/uciml/breast-cancer-wisconsin-data/version/2>) created by Dr. William H. Wolberg, physician at the University Of Wisconsin Hospital at Madison, Wisconsin, USA. The data used for this project was collected in 1993 by the University of Wisconsin and it is composed by the biopsy result of 569 patients in Wisconsin Hospital.

- [Wisconsin Breast Cancer Diagnostic Dataset] <https://www.kaggle.com/uciml/breast-cancer-wisconsin-data/version/2>

The .csv format file containing the data is loaded from my personal github account.

The dataset’s features describe characteristics of the cell nuclei on the image. The features information are specified below:

- Attribute Information:
 1. ID number
 2. Diagnosis (M = malignant, B = benign)
- Ten features were computed for each cell nucleus:
 1. radius: mean of distances from center to points on the perimeter
 2. texture: standard deviation of grey-scale values
 3. perimeter
 4. area: Number of pixels inside contour + $\frac{1}{2}$ for pixels on perimeter
 5. smoothness: local variation in radius lengths), , t
 6. compactness: $\text{perimeter}^2 / \text{area} - 1.0$; this dimensionless number is at a minimum with a circular disk and increases with the irregularity of the boundary, but this measure also increases for elongated cell nuclei, which is not indicative of malignancy
 7. concavity: severity of concave portions of the contour
 8. concave points: number of concave portions of the contour
 9. symmetry

10. fractal dimension: “coastline approximation” - 1; a higher value corresponds a less regular contour and thus to a higher probability of malignancy

The mean, standard error and “worst” or largest (mean of the three largest values) of these features were computed for each image, resulting in 30 variables. From this diagnosis, 357 of the cases were classified as benign tumors and 212 were considered malignant tumors. All cancers and some of the benign masses were histologically confirmed

The column 33 is invalid.

```
data$diagnosis <- as.factor(data$diagnosis)
# the 33 column is invalid
data[,33] <- NULL
```

Chapter 2

Methods and Analysis

2.1 Data Analysis

By observing our dataset, we found that it contains 569 observations with 32 variables.

```
str(data)
```

```
'data.frame':  569 obs. of  32 variables:
 $ id                : int  842302 842517 84300903 84348301 84358402 843786 844359 84458202 844981
 $ diagnosis         : Factor w/ 2 levels "B","M": 2 2 2 2 2 2 2 2 2 ...
 $ radius_mean       : num  18 20.6 19.7 11.4 20.3 ...
 $ texture_mean      : num  10.4 17.8 21.2 20.4 14.3 ...
 $ perimeter_mean    : num  122.8 132.9 130 77.6 135.1 ...
 $ area_mean         : num  1001 1326 1203 386 1297 ...
 $ smoothness_mean   : num  0.1184 0.0847 0.1096 0.1425 0.1003 ...
 $ compactness_mean  : num  0.2776 0.0786 0.1599 0.2839 0.1328 ...
 $ concavity_mean    : num  0.3001 0.0869 0.1974 0.2414 0.198 ...
 $ concave.points_mean : num  0.1471 0.0702 0.1279 0.1052 0.1043 ...
 $ symmetry_mean     : num  0.242 0.181 0.207 0.26 0.181 ...
 $ fractal_dimension_mean : num  0.0787 0.0567 0.06 0.0974 0.0588 ...
 $ radius_se         : num  1.095 0.543 0.746 0.496 0.757 ...
 $ texture_se        : num  0.905 0.734 0.787 1.156 0.781 ...
 $ perimeter_se      : num  8.59 3.4 4.58 3.44 5.44 ...
 $ area_se           : num  153.4 74.1 94 27.2 94.4 ...
 $ smoothness_se     : num  0.0064 0.00522 0.00615 0.00911 0.01149 ...
 $ compactness_se    : num  0.049 0.0131 0.0401 0.0746 0.0246 ...
 $ concavity_se      : num  0.0537 0.0186 0.0383 0.0566 0.0569 ...
 $ concave.points_se : num  0.0159 0.0134 0.0206 0.0187 0.0188 ...
 $ symmetry_se       : num  0.03 0.0139 0.0225 0.0596 0.0176 ...
 $ fractal_dimension_se : num  0.00619 0.00353 0.00457 0.00921 0.00511 ...
 $ radius_worst      : num  25.4 25 23.6 14.9 22.5 ...
 $ texture_worst     : num  17.3 23.4 25.5 26.5 16.7 ...
 $ perimeter_worst   : num  184.6 158.8 152.5 98.9 152.2 ...
 $ area_worst        : num  2019 1956 1709 568 1575 ...
 $ smoothness_worst  : num  0.162 0.124 0.144 0.21 0.137 ...
 $ compactness_worst : num  0.666 0.187 0.424 0.866 0.205 ...
 $ concavity_worst   : num  0.712 0.242 0.45 0.687 0.4 ...
 $ concave.points_worst : num  0.265 0.186 0.243 0.258 0.163 ...
 $ symmetry_worst    : num  0.46 0.275 0.361 0.664 0.236 ...
```

```
$ fractal_dimension_worst: num 0.1189 0.089 0.0876 0.173 0.0768 ...
```

```
head(data)
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
1	842302	M	17.99	10.38	122.80	1001.0
2	842517	M	20.57	17.77	132.90	1326.0
3	84300903	M	19.69	21.25	130.00	1203.0
4	84348301	M	11.42	20.38	77.58	386.1
5	84358402	M	20.29	14.34	135.10	1297.0
6	843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
1	0.11840	0.27760	0.3001	0.14710
2	0.08474	0.07864	0.0869	0.07017
3	0.10960	0.15990	0.1974	0.12790
4	0.14250	0.28390	0.2414	0.10520
5	0.10030	0.13280	0.1980	0.10430
6	0.12780	0.17000	0.1578	0.08089

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
1	0.2419	0.07871	1.0950	0.9053	8.589
2	0.1812	0.05667	0.5435	0.7339	3.398
3	0.2069	0.05999	0.7456	0.7869	4.585
4	0.2597	0.09744	0.4956	1.1560	3.445
5	0.1809	0.05883	0.7572	0.7813	5.438
6	0.2087	0.07613	0.3345	0.8902	2.217

	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
1	153.40	0.006399	0.04904	0.05373	0.01587
2	74.08	0.005225	0.01308	0.01860	0.01340
3	94.03	0.006150	0.04006	0.03832	0.02058
4	27.23	0.009110	0.07458	0.05661	0.01867
5	94.44	0.011490	0.02461	0.05688	0.01885
6	27.19	0.007510	0.03345	0.03672	0.01137

	symmetry_se	fractal_dimension_se	radius_worst	texture_worst
1	0.03003	0.006193	25.38	17.33
2	0.01389	0.003532	24.99	23.41
3	0.02250	0.004571	23.57	25.53
4	0.05963	0.009208	14.91	26.50
5	0.01756	0.005115	22.54	16.67
6	0.02165	0.005082	15.47	23.75

	perimeter_worst	area_worst	smoothness_worst	compactness_worst
1	184.60	2019.0	0.1622	0.6656
2	158.80	1956.0	0.1238	0.1866
3	152.50	1709.0	0.1444	0.4245
4	98.87	567.7	0.2098	0.8663
5	152.20	1575.0	0.1374	0.2050
6	103.40	741.6	0.1791	0.5249

	concavity_worst	concave.points_worst	symmetry_worst
1	0.7119	0.2654	0.4601
2	0.2416	0.1860	0.2750
3	0.4504	0.2430	0.3613
4	0.6869	0.2575	0.6638
5	0.4000	0.1625	0.2364
6	0.5355	0.1741	0.3985

	fractal_dimension_worst
1	0.11890

```

2          0.08902
3          0.08758
4          0.17300
5          0.07678
6          0.12440

```

```
summary(data)
```

```

      id      diagnosis radius_mean texture_mean
Min.   : 8670      B:357   Min.    : 6.981   Min.    : 9.71
1st Qu.: 869218    M:212   1st Qu.:11.700 1st Qu.:16.17
Median : 906024                    Median :13.370 Median :18.84
Mean   : 30371831                Mean   :14.127 Mean   :19.29
3rd Qu.: 8813129                3rd Qu.:15.780 3rd Qu.:21.80
Max.   :911320502                Max.   :28.110 Max.   :39.28

perimeter_mean area_mean smoothness_mean compactness_mean
Min.   : 43.79   Min.   : 143.5   Min.   :0.05263   Min.   :0.01938
1st Qu.: 75.17   1st Qu.: 420.3   1st Qu.:0.08637   1st Qu.:0.06492
Median : 86.24   Median : 551.1   Median :0.09587   Median :0.09263
Mean   : 91.97   Mean   : 654.9   Mean   :0.09636   Mean   :0.10434
3rd Qu.:104.10   3rd Qu.: 782.7   3rd Qu.:0.10530   3rd Qu.:0.13040
Max.   :188.50   Max.   :2501.0   Max.   :0.16340   Max.   :0.34540

concavity_mean concave.points_mean symmetry_mean
Min.   :0.00000   Min.   :0.00000   Min.   :0.1060
1st Qu.:0.02956   1st Qu.:0.02031   1st Qu.:0.1619
Median :0.06154   Median :0.03350   Median :0.1792
Mean   :0.08880   Mean   :0.04892   Mean   :0.1812
3rd Qu.:0.13070   3rd Qu.:0.07400   3rd Qu.:0.1957
Max.   :0.42680   Max.   :0.20120   Max.   :0.3040

fractal_dimension_mean radius_se texture_se perimeter_se
Min.   :0.04996        Min.   :0.1115   Min.   :0.3602   Min.   : 0.757
1st Qu.:0.05770        1st Qu.:0.2324   1st Qu.:0.8339   1st Qu.: 1.606
Median :0.06154        Median :0.3242   Median :1.1080   Median : 2.287
Mean   :0.06280        Mean   :0.4052   Mean   :1.2169   Mean   : 2.866
3rd Qu.:0.06612        3rd Qu.:0.4789   3rd Qu.:1.4740   3rd Qu.: 3.357
Max.   :0.09744        Max.   :2.8730   Max.   :4.8850   Max.   :21.980

area_se smoothness_se compactness_se concavity_se
Min.   : 6.802   Min.   :0.001713   Min.   :0.002252   Min.   :0.00000
1st Qu.: 17.850   1st Qu.:0.005169   1st Qu.:0.013080   1st Qu.:0.01509
Median : 24.530   Median :0.006380   Median :0.020450   Median :0.02589
Mean   : 40.337   Mean   :0.007041   Mean   :0.025478   Mean   :0.03189
3rd Qu.: 45.190   3rd Qu.:0.008146   3rd Qu.:0.032450   3rd Qu.:0.04205
Max.   :542.200   Max.   :0.031130   Max.   :0.135400   Max.   :0.39600

concave.points_se symmetry_se fractal_dimension_se
Min.   :0.000000   Min.   :0.007882   Min.   :0.0008948
1st Qu.:0.007638   1st Qu.:0.015160   1st Qu.:0.0022480
Median :0.010930   Median :0.018730   Median :0.0031870
Mean   :0.011796   Mean   :0.020542   Mean   :0.0037949
3rd Qu.:0.014710   3rd Qu.:0.023480   3rd Qu.:0.0045580
Max.   :0.052790   Max.   :0.078950   Max.   :0.0298400

radius_worst texture_worst perimeter_worst area_worst
Min.   : 7.93   Min.   :12.02   Min.   : 50.41   Min.   : 185.2
1st Qu.:13.01   1st Qu.:21.08   1st Qu.: 84.11   1st Qu.: 515.3
Median :14.97   Median :25.41   Median : 97.66   Median : 686.5
Mean   :16.27   Mean   :25.68   Mean   :107.26   Mean   : 880.6

```


3rd Qu.:18.79	3rd Qu.:29.72	3rd Qu.:125.40	3rd Qu.:1084.0
Max. :36.04	Max. :49.54	Max. :251.20	Max. :4254.0
smoothness_worst	compactness_worst	concavity_worst	concave.points_worst
Min. :0.07117	Min. :0.02729	Min. :0.0000	Min. :0.00000
1st Qu.:0.11660	1st Qu.:0.14720	1st Qu.:0.1145	1st Qu.:0.06493
Median :0.13130	Median :0.21190	Median :0.2267	Median :0.09993
Mean :0.13237	Mean :0.25427	Mean :0.2722	Mean :0.11461
3rd Qu.:0.14600	3rd Qu.:0.33910	3rd Qu.:0.3829	3rd Qu.:0.16140
Max. :0.22260	Max. :1.05800	Max. :1.2520	Max. :0.29100
symmetry_worst	fractal_dimension_worst		
Min. :0.1565	Min. :0.05504		
1st Qu.:0.2504	1st Qu.:0.07146		
Median :0.2822	Median :0.08004		
Mean :0.2901	Mean :0.08395		
3rd Qu.:0.3179	3rd Qu.:0.09208		
Max. :0.6638	Max. :0.20750		

We have to check if the dataset has any missing value:

```
$id
[1] 0
```

```
$diagnosis
[1] 0
```

```
$radius_mean
[1] 0
```

```
$texture_mean
[1] 0
```

```
$perimeter_mean
[1] 0
```

```
$area_mean
[1] 0
```

```
$smoothness_mean
[1] 0
```

```
$compactness_mean
[1] 0
```

```
$concavity_mean
[1] 0
```

```
$concave.points_mean
[1] 0
```

```
$symmetry_mean
[1] 0
```

```
$fractal_dimension_mean
[1] 0
```

```
$radius_se  
[1] 0
```

```
$texture_se  
[1] 0
```

```
$perimeter_se  
[1] 0
```

```
$area_se  
[1] 0
```

```
$smoothness_se  
[1] 0
```

```
$compactness_se  
[1] 0
```

```
$concavity_se  
[1] 0
```

```
$concave.points_se  
[1] 0
```

```
$symmetry_se  
[1] 0
```

```
$fractal_dimension_se  
[1] 0
```

```
$radius_worst  
[1] 0
```

```
$texture_worst  
[1] 0
```

```
$perimeter_worst  
[1] 0
```

```
$area_worst  
[1] 0
```

```
$smoothness_worst  
[1] 0
```

```
$compactness_worst  
[1] 0
```

```
$concavity_worst  
[1] 0
```

```
$concave.points_worst  
[1] 0
```

```
$symmetry_worst
[1] 0
```

```
$fractal_dimension_worst
[1] 0
```

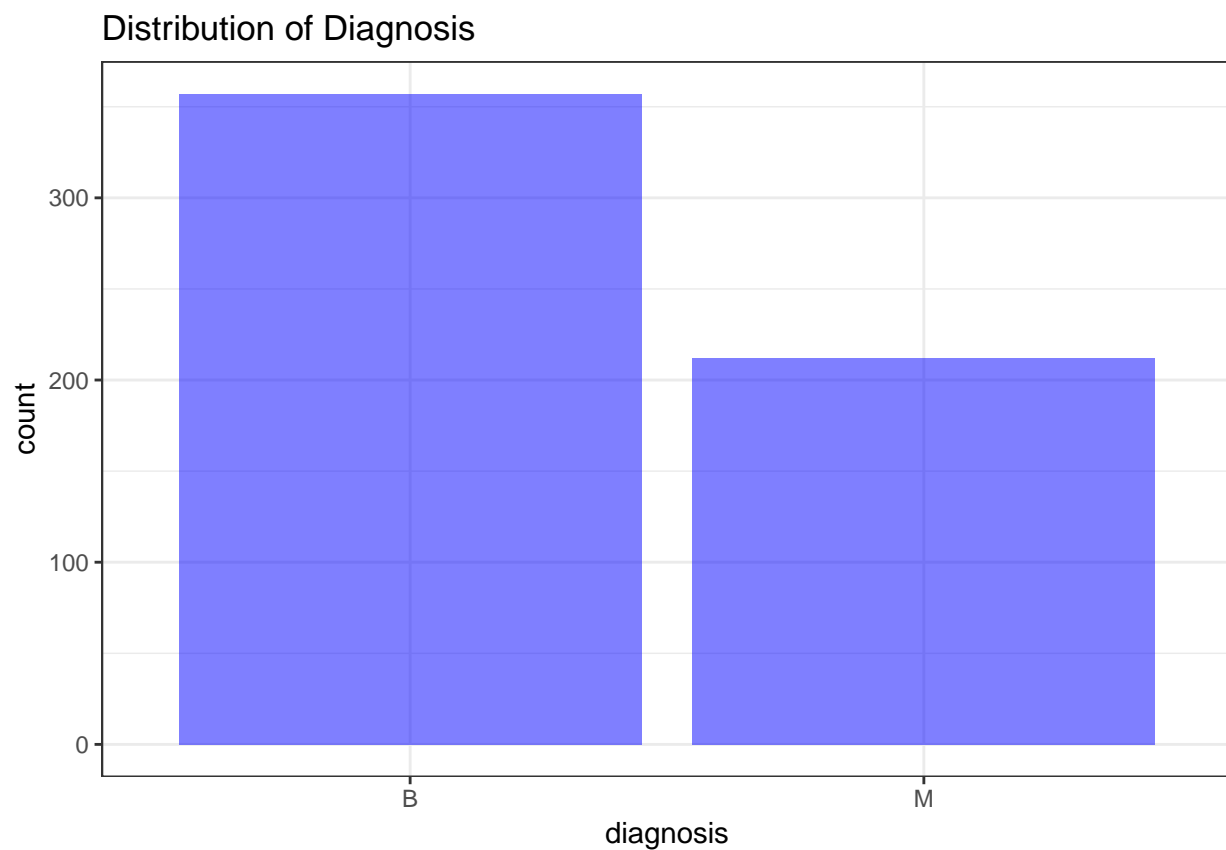
It results that there aren't NA values. By analysing the the dataset we discover that it is a bit unbalanced in its proportions:

```
prop.table(table(data$diagnosis))
```

```
      B      M
0.6274165 0.3725835
```

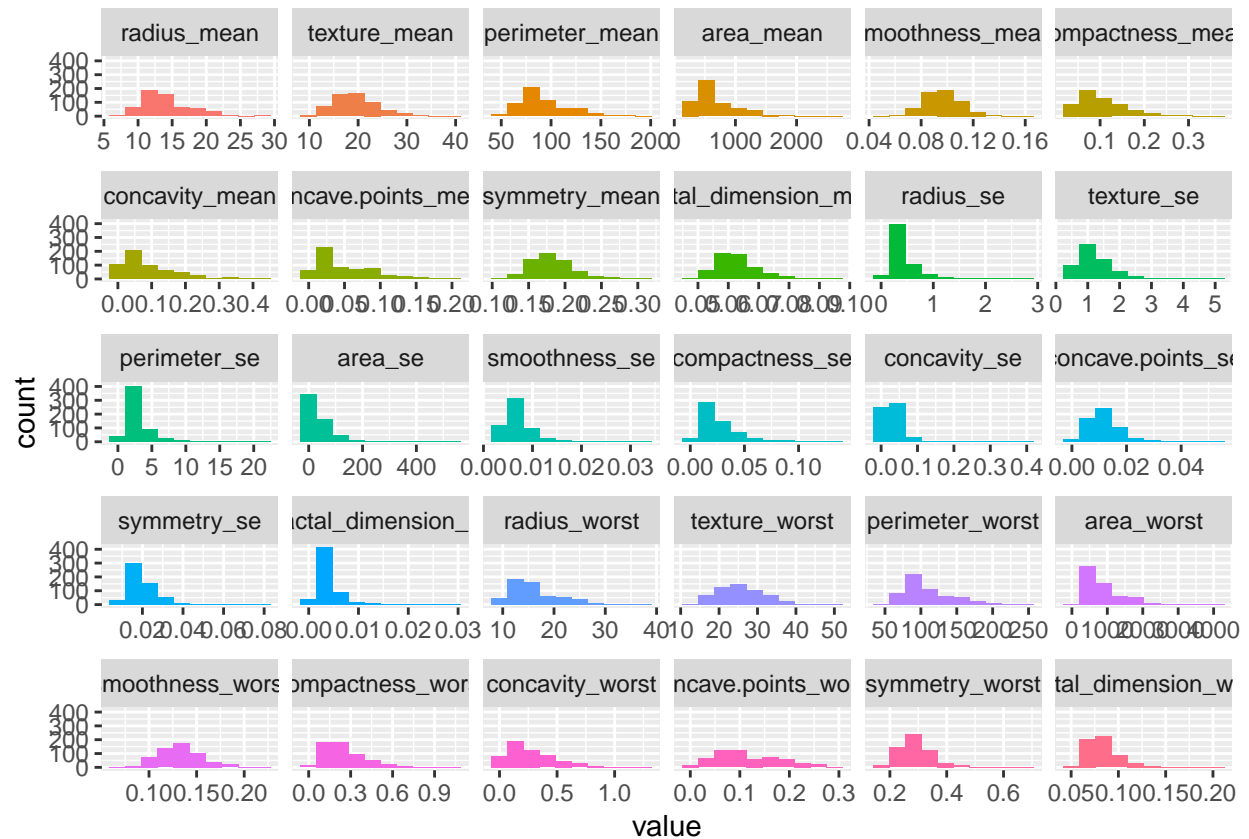
Also the plot of proportions confirms that the target variable is slightly unbalanced.

```
options(repr.plot.width=4, repr.plot.height=4)
ggplot(data, aes(x=diagnosis))+geom_bar(fill="blue",alpha=0.5)+theme_bw()+labs(title="Distribution of D
```




The most variables of the dataset are normally distributed as show with the below plot:

```
plot_num(data %>% select(-id), bins=10)
```



Now we have to check if there is any correlation between variables as machine learning algorithms assume that the predictor variables are independent from each other.

```
correlationMatrix <- cor(data[,3:ncol(data)])
corrplot(correlationMatrix, order = "hclust", tl.cex = 1, addrect = 8)
```



As shown by this plot, many variables are highly correlated with each others. Many methods perform better if highly correlated attributes are removed. The Caret R package provides the `findCorrelation` which will analyze a correlation matrix of your data's attributes report on attributes that can be removed. Because of much correlation some machine learning models could fail.

```
# find attributes that are highly corrected (ideally >0.90)
highlyCorrelated <- findCorrelation(correlationMatrix, cutoff=0.9)
# print indexes of highly correlated attributes
print(highlyCorrelated)
```

```
[1] 7 8 23 21 3 24 1 13 14 2
```

Selecting the right features in our data can mean the difference between mediocre performance with long training times and great performance with short training times.

```
# Remove correlated variables
data2 <- data %>% select(-highlyCorrelated)
# number of columns after removing correlated variables
ncol(data2)
```

```
[1] 22
```

The new dataset has lost 10 variables.


2.2 Modelling Approach

2.2.1 Modelling

Principal Component Analysis (PCA).

To avoid redundancy and relevancy, we used the function ‘prncomp’ to calculate the Principal Component Analysis (PCA) and select the rights components to avoid correlated variables that can be detrimental to our clustering analysis. One of the common problems in analysis of complex data comes from a large number of variables, which requires a large amount of memory and computation power. This is where PCA comes in. It is a technique to reduce the dimension of the feature space by feature extraction. The main idea of PCA is to reduce the dimensionality of a data set consisting of many variables correlated with each other, either heavily or lightly, while retaining the variation present in the dataset, up to the maximum extent. The same is done by transforming the variables to a new set of variables, which are known as the principal components (or simply, the PCs) and are orthogonal, ordered such that the retention of variation present in the original variables decrease as we move down in the order.

```
pca_res_data <- prncomp(data[,3:ncol(data)], center = TRUE, scale = TRUE)
plot(pca_res_data, type="l")
```



```
summary(pca_res_data)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759


	PC7	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.82172	0.69037	0.6457	0.59219	0.5421	0.51104
Proportion of Variance	0.02251	0.01589	0.0139	0.01169	0.0098	0.00871
Cumulative Proportion	0.91010	0.92598	0.9399	0.95157	0.9614	0.97007

	PC13	PC14	PC15	PC16	PC17	PC18
Standard deviation	0.49128	0.39624	0.30681	0.28260	0.24372	0.22939

Proportion of Variance	0.00805	0.00523	0.00314	0.00266	0.00198	0.00175
Cumulative Proportion	0.97812	0.98335	0.98649	0.98915	0.99113	0.99288
	PC19	PC20	PC21	PC22	PC23	PC24
Standard deviation	0.22244	0.17652	0.1731	0.16565	0.15602	0.1344
Proportion of Variance	0.00165	0.00104	0.0010	0.00091	0.00081	0.0006
Cumulative Proportion	0.99453	0.99557	0.9966	0.99749	0.99830	0.9989
	PC25	PC26	PC27	PC28	PC29	PC30
Standard deviation	0.12442	0.09043	0.08307	0.03987	0.02736	0.01153
Proportion of Variance	0.00052	0.00027	0.00023	0.00005	0.00002	0.00000
Cumulative Proportion	0.99942	0.99969	0.99992	0.99997	1.00000	1.00000

As we can observe from the above table, the two first components explains the 0.6324 of the variance. We need 10 principal components to explain more than 0.95 of the variance and 17 to explain more than 0.99.

```
pca_res_data2 <- prcomp(data2[,3:ncol(data2)], center = TRUE, scale = TRUE)
plot(pca_res_data2, type="l")
```



```
summary(pca_res_data2)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.0980	2.0196	1.4663	1.0845	0.91561	0.77019	0.57227
Proportion of Variance	0.4799	0.2039	0.1075	0.0588	0.04192	0.02966	0.01637
Cumulative Proportion	0.4799	0.6838	0.7913	0.8501	0.89205	0.92171	0.93808
	PC8	PC9	PC10	PC11	PC12	PC13	
Standard deviation	0.53641	0.50898	0.45726	0.36641	0.31778	0.28802	
Proportion of Variance	0.01439	0.01295	0.01045	0.00671	0.00505	0.00415	
Cumulative Proportion	0.95247	0.96542	0.97588	0.98259	0.98764	0.99179	
	PC14	PC15	PC16	PC17	PC18	PC19	
Standard deviation	0.21369	0.1846	0.15579	0.15393	0.14782	0.09636	
Proportion of Variance	0.00228	0.0017	0.00121	0.00118	0.00109	0.00046	
Cumulative Proportion	0.99407	0.9958	0.99699	0.99817	0.99926	0.99973	
	PC20						
Standard deviation	0.07375						

```
Proportion of Variance 0.00027
Cumulative Proportion  1.00000
```


The above table shows that 95% of the variance is explained with 8 PC's in the transformed dataset data2.

```
pca_df <- as.data.frame(pca_res_data2$x)
ggplot(pca_df, aes(x=PC1, y=PC2, col=data$diagnosis)) + geom_point(alpha=0.5)
```



The data of the first 2 components can be easily separated into two classes. This is caused by the fact that the variance explained by these components is not large. The data can be easily separated.

```
g_pc1 <- ggplot(pca_df, aes(x=PC1, fill=data$diagnosis)) + geom_density(alpha=0.25)
g_pc2 <- ggplot(pca_df, aes(x=PC2, fill=data$diagnosis)) + geom_density(alpha=0.25)
grid.arrange(g_pc1, g_pc2, ncol=2)
```

Linear Discriminant Analysis (LDA)

Another approach is to use the Linear Discriminant Analysis (LDA) instead of PCA. LDA takes in consideration the different classes and could get better results. The particularity of LDA is that it models the distribution of predictors separately in each of the response classes, and then it uses Bayes' theorem to estimate the probability. It is important to know that LDA assumes a normal distribution for each class, a class-specific mean, and a common variance.

```
lda_res_data <- MASS::lda(diagnosis~., data = data, center = TRUE, scale = TRUE)
lda_res_data
```

Call:

```
lda(diagnosis ~ ., data = data, center = TRUE, scale = TRUE)
```

Prior probabilities of groups:

	B	M
Prior probabilities	0.6274165	0.3725835

Group means:

	id	radius_mean	texture_mean	perimeter_mean	area_mean
B	26543825	12.14652	17.91476	78.07541	462.7902
M	36818050	17.46283	21.60491	115.36538	978.3764

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
B	0.09247765	0.08008462	0.04605762	0.02571741
M	0.10289849	0.14518778	0.16077472	0.08799000

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
B	0.174186	0.06286739	0.2840824	1.220380	2.000321
M	0.192909	0.06268009	0.6090825	1.210915	4.323929

```

      area_se smoothness_se compactness_se concavity_se concave.points_se
B 21.13515   0.007195902     0.02143825   0.02599674     0.009857653
M 72.67241   0.006780094     0.03228117   0.04182401     0.015060472

      symmetry_se fractal_dimension_se radius_worst texture_worst
B 0.02058381     0.003636051     13.37980     23.51507
M 0.02047240     0.004062406     21.13481     29.31821

      perimeter_worst area_worst smoothness_worst compactness_worst
B      87.00594     558.8994     0.1249595     0.1826725
M     141.37033    1422.2863     0.1448452     0.3748241

      concavity_worst concave.points_worst symmetry_worst
B      0.1662377     0.07444434     0.2702459
M      0.4506056     0.18223731     0.3234679

      fractal_dimension_worst
B              0.07944207
M              0.09152995

```

Coefficients of linear discriminants:

```


                                LD1
id                             -2.512117e-10
radius_mean                    -1.080876e+00
texture_mean                    2.338408e-02
perimeter_mean                 1.172707e-01
area_mean                      1.595690e-03
smoothness_mean                5.251575e-01
compactness_mean               -2.094197e+01
concavity_mean                 6.955923e+00
concave.points_mean            1.047567e+01
symmetry_mean                  4.938898e-01
fractal_dimension_mean         -5.937663e-02
radius_se                      2.101503e+00
texture_se                     -3.979869e-02
perimeter_se                   -1.121814e-01
area_se                        -4.083504e-03
smoothness_se                  7.987663e+01
compactness_se                 1.387026e-01
concavity_se                   -1.768261e+01
concave.points_se              5.350520e+01
symmetry_se                    8.143611e+00
fractal_dimension_se           -3.431356e+01
radius_worst                   9.677207e-01
texture_worst                  3.540591e-02
perimeter_worst                -1.204507e-02
area_worst                     -5.012127e-03
smoothness_worst               2.612258e+00
compactness_worst              3.636892e-01
concavity_worst                1.880699e+00
concave.points_worst           2.218189e+00
symmetry_worst                 2.783102e+00
fractal_dimension_worst        2.117830e+01

```

#Data frame of the LDA for visualization purposes

```
lda_df_predict <- predict(lda_res_data, data)$x %>% as.data.frame() %>% cbind(diagnosis=data$diagnosis)
```

```
ggplot(lda_df_predict, aes(x=LD1, fill=diagnosis)) + geom_density(alpha=0.5)
```



2.2.2 Model creation

We are going to get a training and a testing set to use when building some models. We split the modified dataset into Train (80%) and Test (20%), in order to predict is whether a cancer cell is Benign or Malignant, by building machine learning classification models.

```
set.seed(1815)
data3 <- cbind (diagnosis=data$diagnosis, data2)
data_sampling_index <- createDataPartition(data$diagnosis, times=1, p=0.8, list = FALSE)
train_data <- data3[data_sampling_index, ]
test_data <- data3[-data_sampling_index, ]

fitControl <- trainControl(method="cv",      #Control the computational nuances of the train function
                           number = 15,    #Either the number of folds or number of resampling iterations
                           classProbs = TRUE,
                           summaryFunction = twoClassSummary)
```

2.2.3 Naive Bayes Model

The Naive Bayesian classifier is based on Bayes' theorem with the independence assumptions between predictors. A Naive Bayesian model is easy to build, with no complicated iterative parameter estimation which makes it particularly useful for very large datasets. Bayes theorem provides a way of calculating the posterior probability, $P(c|x)$, from $P(c)$, $P(x)$, and $P(x|c)$. Naive Bayes classifier assume that the effect of the

value of a predictor (x) on a given class (c) is independent of the values of other predictors. This assumption is called class conditional independence.

```
model_naiveb <- train(diagnosis~.,
                      train_data,
                      method="nb",
                      metric="ROC",
                      preProcess=c('center', 'scale'), #in order to normalize the data
                      trace=FALSE,
                      trControl=fitControl)

prediction_naiveb <- predict(model_naiveb, test_data)
confusionmatrix_naiveb <- confusionMatrix(prediction_naiveb, test_data$diagnosis, positive = "M")
confusionmatrix_naiveb
```

Confusion Matrix and Statistics

	Reference	
Prediction	B	M
B	69	5
M	2	37

Accuracy : 0.9381
 95% CI : (0.8765, 0.9747)
 No Information Rate : 0.6283
 P-Value [Acc > NIR] : 1.718e-14

 Kappa : 0.8654
 McNemar's Test P-Value : 0.4497


 Sensitivity : 0.8810
 Specificity : 0.9718
 Pos Pred Value : 0.9487
 Neg Pred Value : 0.9324
 Prevalence : 0.3717
 Detection Rate : 0.3274
 Detection Prevalence : 0.3451
 Balanced Accuracy : 0.9264

 'Positive' Class : M

We can note the accuracy with such model. We will later describe better these metrics, where: Sensitivity (recall) represent the true positive rate: the proportions of actual positives correctly identified. Specificity is the true negative rate: the proportion of actual negatives correctly identified. Accuracy is the general score of the classifier model performance as it is the ratio of how many samples are correctly classified to all samples. F1 score: the harmonic mean of precision and sensitivity. Accuracy and F1 score would be used to compare the result with the benchmark model. Precision: the number of correct positive results divided by the number of all positive results returned by the classifier.

The most important variables that permit the best prediction and contribute the most to the model are the following:

```
plot(varImp(model_naiveb), top=10, main="Top variables- Naive Bayes")
```



2.2.4 Logistic Regression Model

Logistic Regression is widely used for binary classification like (0,1). The binary logistic model is used to estimate the probability of a binary response based on one or more predictor (or independent) variables (features).

```
model_logreg<- train(diagnosis ~., data = train_data, method = "glm",
                     metric = "ROC",

                     preProcess = c("scale", "center"), # in order to normalize the data
                     trControl= fitControl)
prediction_logreg<- predict(model_logreg, test_data)

# Check results
confusionmatrix_logreg <- confusionMatrix(prediction_logreg, test_data$diagnosis, positive = "M")
confusionmatrix_logreg
```

Confusion Matrix and Statistics

	Reference	
Prediction	B	M
B	71	2
M	0	40

Accuracy : 0.9823

```

          95% CI : (0.9375, 0.9978)
No Information Rate : 0.6283
P-Value [Acc > NIR] : <2e-16

          Kappa : 0.9617
McNemar's Test P-Value : 0.4795


          Sensitivity : 0.9524
          Specificity : 1.0000
          Pos Pred Value : 1.0000
          Neg Pred Value : 0.9726
          Prevalence : 0.3717
          Detection Rate : 0.3540
          Detection Prevalence : 0.3540
          Balanced Accuracy : 0.9762

          'Positive' Class : M

```

The most important variables that permit the best prediction and contribute the most to the model are the following:

```
plot(varImp(model_logreg), top=10, main="Top variables - Log Regr")
```



2.2.5 Random Forest Model

Random forests are a very popular machine learning approach that addresses the shortcomings of decision trees using a clever idea. The goal is to improve prediction performance and reduce instability by averaging multiple decision trees (a forest of trees constructed with randomness). Random forest is another ensemble method based on decision trees. It split data into sub-samples, trains decision tree classifiers on each sub-sample and averages prediction of each classifier. Splitting dataset causes higher bias but it is compensated by large decrease in variance. Random Forest is a supervised learning algorithm and it is flexible, easy to use machine learning algorithm that produces, even without hyper-parameter tuning, a great result most of the time. It is also one of the most used algorithms, because of it's simplicity and the fact that it can be used for both classification and regression tasks. Random forest builds multiple decision trees and merges them together to get a more accurate and stable prediction.

```
model_randomforest <- train(diagnosis~.,
                           train_data,
                           method="rf", #also recommended ranger, because it is a lot faster than ori,
                           metric="ROC",
                           #tuneLength=10,
                           #tuneGrid = expand.grid(mtry = c(2, 3, 6)),
                           preProcess = c('center', 'scale'),
                           trControl=fitControl)

prediction_randomforest <- predict(model_randomforest, test_data)

#Check results
confusionmatrix_randomforest <- confusionMatrix(prediction_randomforest, test_data$diagnosis, positive = "M")
confusionmatrix_randomforest
```

Confusion Matrix and Statistics

		Reference	
Prediction		B	M
B	71	3	
M	0	39	


Accuracy : 0.9735
 95% CI : (0.9244, 0.9945)
 No Information Rate : 0.6283
 P-Value [Acc > NIR] : <2e-16

 Kappa : 0.9423
 McNemar's Test P-Value : 0.2482

 Sensitivity : 0.9286
 Specificity : 1.0000
 Pos Pred Value : 1.0000
 Neg Pred Value : 0.9595
 Prevalence : 0.3717
 Detection Rate : 0.3451
 Detection Prevalence : 0.3451
 Balanced Accuracy : 0.9643

 'Positive' Class : M

```
plot(varImp(model_randomforest), top=10, main="Top variables- Random Forest")
```



2.2.6 K Nearest Neighbor (KNN) Model

KNN (K-Nearest Neighbors) is one of many (supervised learning) algorithms used in data mining and machine learning, it's a classifier algorithm where the learning is based "how similar" is a data from other. K nearest neighbors is a simple algorithm that stores all available cases and classifies new cases based on a similarity measure (e.g., distance functions).

```
model_knn <- train(diagnosis~.,
  train_data,
  method="knn",
  metric="ROC",
  preProcess = c('center', 'scale'),
  tuneLength=10, #The tuneLength parameter tells the algorithm to try different default
  #In this case we used 10 default values
  trControl=fitControl)

prediction_knn <- predict(model_knn, test_data)
confusionmatrix_knn <- confusionMatrix(prediction_knn, test_data$diagnosis, positive = "M")
confusionmatrix_knn
```

Confusion Matrix and Statistics

	Reference	
Prediction	B	M
B	70	5


```

M 1 37

Accuracy : 0.9469
 95% CI : (0.888, 0.9803)
No Information Rate : 0.6283
P-Value [Acc > NIR] : 1.866e-15

Kappa : 0.8841
McNemar's Test P-Value : 0.2207


Sensitivity : 0.8810
Specificity : 0.9859
Pos Pred Value : 0.9737
Neg Pred Value : 0.9333
Prevalence : 0.3717
Detection Rate : 0.3274
Detection Prevalence : 0.3363
Balanced Accuracy : 0.9334

'Positive' Class : M

```

The most important variables that permit the best prediction and contribute the most to the model are the following:

```
plot(varImp(model_knn), top=10, main="Top variables - KNN")
```



2.2.7 Neural Network with PCA Model

Artificial Neural Networks (NN) are a types of mathematical algorithms originating in the simulation of networks of biological neurons. An artificial Neural Network consists of nodes (called neurons) and edges (called synapses). Input data is transmitted through the weighted synapses to the neurons where calculations are processed and then either sent to further neurons or represent the output.

Neural Networks take in the weights of connections between neurons . The weights are balanced, learning data point in the wake of learning data point . When all weights are trained, the neural network can be utilized to predict the class or a quantity, if there should arise an occurrence of regression of a new input data point. With Neural networks, extremely complex models can be trained and they can be utilized as a kind of black box, without playing out an unpredictable complex feature engineering before training the model. Joined with the “deep approach” even more unpredictable models can be picked up to realize new possibilities.

```
model_nnet_pca <- train(diagnosis~.,
                        train_data,
                        method="nnet",
                        metric="ROC",
                        preProcess=c('center', 'scale', 'pca'),
                        tuneLength=10,
                        trace=FALSE,
                        trControl=fitControl)

prediction_nnet_pca <- predict(model_nnet_pca, test_data)
confusionmatrix_nnet_pca <- confusionMatrix(prediction_nnet_pca, test_data$diagnosis, positive = "M")
confusionmatrix_nnet_pca
```

Confusion Matrix and Statistics

```

      Reference
Prediction B  M
      B  71  2
      M   0 40

      Accuracy : 0.9823
      95% CI : (0.9375, 0.9978)
      No Information Rate : 0.6283
      P-Value [Acc > NIR] : <2e-16

      Kappa : 0.9617
      Mcnemar's Test P-Value : 0.4795


      Sensitivity : 0.9524
      Specificity : 1.0000
      Pos Pred Value : 1.0000
      Neg Pred Value : 0.9726
      Prevalence : 0.3717
      Detection Rate : 0.3540
      Detection Prevalence : 0.3540
      Balanced Accuracy : 0.9762

      'Positive' Class : M
```

The most important variables that permit the best prediction and contribute the most to the model are the

following:

```
plot(varImp(model_nnet_pca), top=8, main="Top variables - NNET PCA")
```



2.2.8 Neural Network with LDA Model

We are going to create a training and test set of LDA data created in previous chapters:

```
train_data_lda <- lda_df_predict[data_sampling_index, ]
test_data_lda <- lda_df_predict[-data_sampling_index, ]
```

```
model_nnet_lda <- train(diagnosis~.,
                        train_data_lda,
                        method="nnet",
                        metric="ROC",
                        preProcess=c('center', 'scale'),
                        tuneLength=10,
                        trace=FALSE,
                        trControl=fitControl)
```

```
prediction_nnet_lda <- predict(model_nnet_lda, test_data_lda)
```

```
confusionmatrix_nnet_lda <- confusionMatrix(prediction_nnet_lda, test_data_lda$diagnosis, positive = "M")
confusionmatrix_nnet_lda
```

Confusion Matrix and Statistics

	Reference	
Prediction	B	M

B 71 1
M 0 41

Accuracy : 0.9912
95% CI : (0.9517, 0.9998)
No Information Rate : 0.6283
P-Value [Acc > NIR] : <2e-16

Kappa : 0.981
McNemar's Test P-Value : 1

Sensitivity : 0.9762
Specificity : 1.0000
Pos Pred Value : 1.0000
Neg Pred Value : 0.9861
Prevalence : 0.3717
Detection Rate : 0.3628
Detection Prevalence : 0.3628
Balanced Accuracy : 0.9881

'Positive' Class : M

Chapter 3

Results

We can now compare and evaluate the results obtained with the above calculations.

```
models_list <- list(Naive_Bayes=model_naiveb,
                    Logistic_regr=model_logreg,
                    Random_Forest=model_randomforest,
                    KNN=model_knn,
                    Neural_PCA=model_nnet_pca,
                    Neural_LDA=model_nnet_lda)
models_results <- resamples(models_list)

summary(models_results)
```

Call:

```
summary.resamples(object = models_results)
```

Models: Naive_Bayes, Logistic_regr, Random_Forest, KNN, Neural_PCA, Neural_LDA
Number of resamples: 15

ROC

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
Naive_Bayes	0.9282297	0.9794657	0.9952153	0.9863636	1	1	0
Logistic_regr	0.8205742	0.9385965	1.0000000	0.9603535	1	1	0
Random_Forest	0.9425837	0.9880383	1.0000000	0.9916680	1	1	0
KNN	0.9188596	0.9813596	1.0000000	0.9885965	1	1	0
Neural_PCA	0.9760766	1.0000000	1.0000000	0.9977671	1	1	0
Neural_LDA	0.9712919	0.9978070	1.0000000	0.9965178	1	1	0

Sens


	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
Naive_Bayes	0.8421053	0.9210526	0.9473684	0.9508772	1	1	0
Logistic_regr	0.8421053	0.9473684	0.9473684	0.9508772	1	1	0
Random_Forest	0.8947368	0.9473684	1.0000000	0.9719298	1	1	0
KNN	0.9473684	1.0000000	1.0000000	0.9964912	1	1	0
Neural_PCA	0.9473684	0.9736842	1.0000000	0.9859649	1	1	0
Neural_LDA	0.8947368	0.9736842	1.0000000	0.9824561	1	1	0

Spec

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
Naive_Bayes	0.7272727	0.8257576	0.9090909	0.8994949	1	1	0
Logistic_regr	0.7500000	0.8712121	1.0000000	0.9419192	1	1	0
Random_Forest	0.6363636	0.8636364	0.9166667	0.9116162	1	1	0
KNN	0.7272727	0.9090909	0.9166667	0.9131313	1	1	0
Neural_PCA	0.8181818	0.9090909	1.0000000	0.9525253	1	1	0
Neural_LDA	0.7272727	0.9090909	1.0000000	0.9520202	1	1	0

As we can observe from the following plot, two models, Naive_bayes and Logistic_regr have great variability, depending of the processed sample :

```
bwplot(models_results, metric="ROC")
```



The Neural Network LDA model achieve a great auc (Area Under the ROC Curve) with some variability. The ROC (Receiver Operating characteristic Curve is a graph showing the performance of a classification model at all classification thresholds) metric measure the auc of the roc curve of each model. This metric is independent of any threshold. Let's remember how these models result with the testing dataset. Prediction classes are obtained by default with a threshold of 0.5 which could not be the best with an unbalanced dataset like this.

```
confusionmatrix_list <- list(
  Naive_Bayes=confusionmatrix_naiveb,
  Logistic_regr=confusionmatrix_logreg,
  Random_Forest=confusionmatrix_randomforest,
  KNN=confusionmatrix_knn,
  Neural_PCA=confusionmatrix_nnet_pca,
  Neural_LDA=confusionmatrix_nnet_lda)
confusionmatrix_list_results <- sapply(confusionmatrix_list, function(x) x$byClass)
confusionmatrix_list_results %>% knitr::kable()
```

	Naive_Bayes	Logistic_regr	Random_Forest	KNN	Neural_PCA	Neural_LDA
Sensitivity	0.8809524	0.9523810	0.9285714	0.8809524	0.9523810	0.9761905
Specificity	0.9718310	1.0000000	1.0000000	0.9859155	1.0000000	1.0000000
Pos Pred Value	0.9487179	1.0000000	1.0000000	0.9736842	1.0000000	1.0000000
Neg Pred Value	0.9324324	0.9726027	0.9594595	0.9333333	0.9726027	0.9861111
Precision	0.9487179	1.0000000	1.0000000	0.9736842	1.0000000	1.0000000
Recall	0.8809524	0.9523810	0.9285714	0.8809524	0.9523810	0.9761905
F1	0.9135802	0.9756098	0.9629630	0.9250000	0.9756098	0.9879518
Prevalence	0.3716814	0.3716814	0.3716814	0.3716814	0.3716814	0.3716814
Detection Rate	0.3274336	0.3539823	0.3451327	0.3274336	0.3539823	0.3628319
Detection Prevalence	0.3451327	0.3539823	0.3451327	0.3362832	0.3539823	0.3628319
Balanced Accuracy	0.9263917	0.9761905	0.9642857	0.9334339	0.9761905	0.9880952

Chapter 4

Discussion

We will now describe the metrics that we will compare in this section.

Accuracy is our starting point. It is the number of correct predictions made divided by the total number of predictions made, multiplied by 100 to turn it into a percentage.

Precision is the number of True Positives divided by the number of True Positives and False Positives. Put another way, it is the number of positive predictions divided by the total number of positive class values predicted. It is also called the Positive Predictive Value (PPV). A low precision can also indicate a large number of False Positives.

Recall (Sensitivity) is the number of True Positives divided by the number of True Positives and the number of False Negatives. Put another way it is the number of positive predictions divided by the number of positive class values in the test data. It is also called Sensitivity or the True Positive Rate. Recall can be thought of as a measure of a classifiers completeness. A low recall indicates many False Negatives.

The F1 Score is the $2 \times ((\text{precision} \times \text{recall}) / (\text{precision} + \text{recall}))$. It is also called the F Score or the F Measure. Put another way, the F1 score conveys the balance between the precision and the recall.

The best results for sensitivity (detection of breast cancer malign cases) is Neural Network with LDA model which also has a great F1 score.

```
confusionmatrix_results_max <- apply(confusionmatrix_list_results, 1, which.is.max)

output_report <- data.frame(metric=names(confusionmatrix_results_max),
                             best_model=colnames(confusionmatrix_list_results)[confusionmatrix_results_max],
                             value=mapapply(function(x,y) {confusionmatrix_list_results[x,y]},
                                              names(confusionmatrix_results_max),
                                              confusionmatrix_results_max))

rownames(output_report) <- NULL
output_report
```

	metric	best_model	value
1	Sensitivity	Neural_LDA	0.9761905
2	Specificity	Logistic_regr	1.0000000
3	Pos Pred Value	Neural_LDA	1.0000000
4	Neg Pred Value	Neural_LDA	0.9861111
5	Precision	Logistic_regr	1.0000000
6	Recall	Neural_LDA	0.9761905
7	F1	Neural_LDA	0.9879518
8	Prevalence	Logistic_regr	0.3716814
9	Detection Rate	Neural_LDA	0.3628319

10	Detection Prevalence	Neural_LDA	0.3628319
11	Balanced Accuracy	Neural_LDA	0.9880952

Chapter 5

Conclusion

This paper treats the Wisconsin Madison Breast Cancer diagnosis problem as a pattern classification problem. In this report we investigated several machine learning model and we selected the optimal model by selecting a high accuracy level combined with a low rate of false-negatives (the means that the metric is high sensitivity).

The Neural Network with LDA model had the optimal results for F1 (0.9879518), Sensitivity (0.9761905) and Balanced Accuracy (0.9880952)

Chapter 6

Appendix - Environment

```
print("Operating System:")
```

```
[1] "Operating System:"
```

```
version
```

```
platform      _  
arch          x86_64  
os            darwin15.6.0  
system        x86_64, darwin15.6.0  
status  
major         3  
minor         5.1  
year          2018  
month         07  
day           02  
svn rev       74947  
language      R  
version.string R version 3.5.1 (2018-07-02)  
nickname      Feather Spray
```